



HOME BIG DATA HADOOP MAP REDUCE HIVE PIG HBASE FLUME INTERVIEW QUESTIONS
MISCELLANEOUS OTHERTOOLS HADOOP DISCUSSION FORUM

Home » Hadoop Common » Hive » Hive Functions Examples



Hive Functions Examples



This entry was posted in [Hive](#) on May 28, 2015 by Siva

Table of Contents [hide]

Hive Functions Examples

- SET
- SHOW
- USE
- CREATE DATABASE
- CREATE MANAGED TABLE
- CREATE EXTERNAL TABLE:
- CREATING TABLE FROM EXISTING TABLE
- CREATING EXTERNAL TABLES FROM MANAGED TABLES:
- LOAD
- COPY DATA FROM ONE TABLE TO ANOTHER
- DROP
- QUIT
- SELECT
- DESCRIBE
- DESCRIBE SPECIFIC FIELD
- DESCRIBE EXTENDED
- ALTER
- CLONE SCHEMA (DATA IS NOT COPIED)
- CLONE SCHEMA TO ANOTHER DB
- USING REGULAR EXPRESSIONS:
- MATHEMATICAL FUNCTIONS:
- AGGREGATE FUNCTIONS:
- LIMIT
- NESTED SELECT STATEMENT
- CASE..WHEN..THEN
- LIKE & RLIKE
- JOINS
- ORDER BY, SORT BY, DISTRIBUTED BY WITH SORT BY, CLUSTER BY
- UNION ALL (combines 2 or more tables)
- CREATE TABLES FROM VIEWS:
- ALTER VIEW:

Popular Pages

[Hive Archives - Hadoop Online Tutorials](#)

[Partitioning in Hive - Hadoop Online Tutorials](#)

More Info

Core Hadoop

-DROP VIEW:
 -INDEX:
 -> Inefficient:
 -> Efficient:
 -INDEXES:
 -EXPLAIN:
 ABSTRACT SYNTAX TREE:
 STAGE DEPENDENCIES:
 STAGE PLANS:
 SCRIPTS:
 -CHANGING HIVE PROMPT:
 -PRINT column headers;
 -HIVE ONESHOT COMMANDS:
 -EXECUTE UNIX COMMANDS:
 -INSTALLATION LOCATION:
 -HIVE CONFIGURATION:
 -HIVE LOG:
 -HIVE LOGGING:
 -HIVE SERVICES:
 -CLI SERVICE
 - HIVE NAMESPACES:
 -SELECT in NON-INTERACTIVE MODE:
 -HIVE FILES:
 -CHANGE TO ROOT USER
 -EXECUTE HADOOP COMMANDS:
 -LIST ALL HADOOP COMMANDS:
 -COMMENTS IN HIVE:
 WORK:
 Share this:

Hive Functions Examples

SET
 SHOW
 USE
 CREATE DATABASE
 CREATE MANAGED TABLE
 CREATE EXTERNAL TABLE
 CREATING TABLE FROM EXISTING TABLE
 CREATING EXTERNAL TABLES FROM MANAGED TABLES
 LOAD
 COPY DATA FROM ONE TABLE TO ANOTHER
 DROP
 QUIT
 SELECT
 DESCRIBE
 DESCRIBE SPECIFIC FIELD
 DESCRIBE EXTENDED

Big Data
 Hadoop
 Map Reduce

EcoSystem Tools

Hive
 Pig
 HBase
 Impala

Sponsored Content



Watch Silicon Valley Free For 30 Days: Click Here
 Hotstar



Top 5: New Luxury SUVs
 Departures International



5 ways to avoid risks in trading
www.5paisa.com



The BMW 3 Series. Celebrate as you beat price hike. Know more.
 BMW India

Recommended by

ALTER
CLONE SCHEMA (DATA IS NOT COPIED)
CLONE SCHEMA TO ANOTHER DB
USING REGULAR EXPRESSIONS
MATHEMATICAL FUNCTIONS
AGGREGATE FUNCTIONS
LIMIT
NESTED SELECT STATEMENT
CASE..WHEN..THEN
LIKE & RLIKE

JOINS

ORDER BY, SORT BY, DISTRIBUTED BY WITH SORT BY, CLUSTER BY

SET

```
hive> set;      --> prints all variables from these 4 namespaces:
```

SHOW

```
hive> show databases;  
hive> show tables;  
hive> show databases like 'f*';
```

USE

```
hive> use default;  
hive (default)> use test;
```

CREATE DATABASE

```
hive> create database if not exists financials;  
hive> create database financials  
> comment 'holds all financial tables'  
> location '/tmp/hive_db_dir';
```

CREATE MANAGED TABLE

```
hive> create table records (year string, temperature int, quantity int)  
> row format delimited  
> fields terminated by '\t';  
hive> create table employees (  
> name string,  
> salary float,  
> subordinates array<string>,  
> deductions map<string, float>,  
> address struct<street:string, city:string, state:string, zip:int>);  
hive> create database financials2  
> with dbproperties('creator' = 'Ganesh Pillai', 'date' = '2015-02-19');
```

CREATE EXTERNAL TABLE:

```
[cloudera@localhost ~]$ cat scripts/create_external_table.hql;  
create external table if not exists mydb.stocks(  
exchange string,
```

```
symbol string,
ymd string,
price_open float,
price_high float,
price_low float,
price_close float,
volume int,
price_adj_close float)
row format delimited fields terminated by ',' location '/user/cloudera/myhive';
```

CREATING TABLE FROM EXISTING TABLE

```
hive (mydb)> create external table stocks2
> like stocks;
```

CREATING EXTERNAL TABLES FROM MANAGED TABLES:

```
hive> create external table mydb.employees3
> like employees;
```

LOAD

```
hive> load data local inpath 'data/sample.txt'
> overwrite into table records;
hive (mydb)> load data local inpath 'data/NYSE_daily'
> overwrite into table stocks;
```

COPY DATA FROM ONE TABLE TO ANOTHER

```
Hive> insert overwrite table nyse_daily_bak select * from nyse_daily;
NOTE: Lauges Map-only job
hive> insert overwrite dividends_bak select * from dividends;
```

DROP

```
hive> drop table records;
hive (financials2)> drop database if exists financials2 cascade;
--NOTE: cascade/restrict is optional
```

QUIT

```
hive> quit;
```

SELECT

```
hive> select year, MAX(temperature) from records
> group by year;
```

DESCRIBE

```
hive> describe database financials;
hive> describe database default;
hive (mydb)> describe database mydb;
hive (mydb)> describe employees;
hive (mydb)> describe extended employees;
hive (mydb)> describe formatted employees;
```

DESCRIBE SPECIFIC FIELD

```
hive (mydb)> describe employees.address;
```

DESCRIBE EXTENDED

```
hive> describe database extended financials2;
OK
financials2          hdfs://localhost.localdomain:8020/user/hive/warehouse
Time taken: 0.124 seconds
```

ALTER

```
hive (financials2)> alter database financials set dbproperties('modified by' = 'GANESH PILLAI');
```

CLONE SCHEMA (DATA IS NOT COPIED)

```
hive (mydb)> create table mydb.employees2
> like employees;
```

CLONE SCHEMA TO ANOTHER DB

```
hive (mydb)> create table financials.employees3 like employees;
hive (mydb)> load data local inpath 'data/NYSE_daily' into table stocks2;
hive> drop database mydb;
FAILED: InvalidOperationException(message:Database mydb is not empty)
hive> drop database mydb cascade;
[cloudera@localhost ~]$ cat scripts/stocks.hql;
create external table if not exists mydb.stocks(exchange string,
symbol string,
ymd string,
price_open float,
price_high float,
price_low float,
price_close float,
volume int,
price_adj_close float)
row format delimited fields terminated by '\t' location '/user/cloudera/myhive';
[cloudera@localhost ~]$ cat scripts/load_stocks.hql;
load data local inpath 'data/NYSE_daily'
overwrite into table mydb.stocks;
```

–EMPLOYEES TABLE:

```
[cloudera@localhost ~]$ cat scripts/employees_no_partition.hql;
create table mydb.employees_no_partition(name string,
salary float,
subordinates array<string>,
deductions map<string, float>,
address struct<street:string, city:string, state:string, zip:int>)
row format delimited fields terminated by '\001'
collection items terminated by '\002'
map keys terminated by '\003'
lines terminated by '\n' stored as textfile;

[cloudera@localhost ~]$ cat scripts/load_employees_no_partition.hql;
load data local inpath 'data/employees.txt'
overwrite into table mydb.employees_no_partition;

hive (mydb)> select name, subordinates[0] from employees_no_partition;
hive (mydb)> select name, address from employees_no_partition;
name                                address
John Doe                            {"street":"1 Michigan Ave.", "city":"Chicago", "state":
hive (mydb)> select name, deductions["State Taxes"] from employees_no_partition;
name                                _c1
John Doe                            0.05

hive (mydb)> select name, address.city from employees_no_partition;
name                                city
John Doe                            Chicago
```

USING REGULAR EXPRESSIONS:

```
hive (mydb)> select symbol, `price.*` from stocks limit 5;
symbol  price_open  price_high  price_low  price_close
CLI     35.39       35.7        34.5
```

COMPUTING WITH COLUMN VALUES(using functions & arithmetic expressions)

```
hive (mydb)> select upper(name), salary, deductions["Federal Taxes"], round(salary*(1-deductio
_c0      salary      _c2      _c3
JOHN DOE      100000.0      0.2      80000
hive (mydb)> select count(*) as count, avg(salary) as salary from employees_no_partition;
count      salary
7           102857.14285714286
```

ARITHMETIC OPERATORS:

+, -, *, /, %, &(AND), |(OR), ^(XOR), ~(NOT)

MATHEMATICAL FUNCTIONS:

round, rand, abs, etc

```
The Numerical functions are listed below in alphabetical order. Use these functions in SQL
ABS( double n )

The ABS function returns the absolute value of a number.

Example: ABS(-100)

ACOS( double n )
```

The ACOS function returns the arc cosine of value n. This function returns Null if the value n is not in the range [-1,1].

Example: ACOS(0.5)

ASIN(double n)

The ASIN function returns the arc sin of value n. This function returns Null if the value n is not in the range [-1,1].

Example: ASIN(0.5)

BIN(bigint n)

The BIN function returns the number n in the binary format.

Example: BIN(100)

CEIL(double n), CEILING(double n)

The CEILING or CEILING function returns the smallest integer greater than or equal to the value n.

Example: CEIL(9.5)

CONV(bigint n, int from_base, int to_base)

The CONV function converts the given number n from one base to another base.

Example: CONV(100, 10,2)

COS(double n)

AGGREGATE FUNCTIONS:

count, sum, ave, min, max, etc

```
hive (mydb)> select count(*), avg(salary) from employees_no_partition;
hive (mydb)> set hive.map.aggr = true;
hive (mydb)> select MAX(column_name) from tablename;
hive (mydb)> select MIN(column_name) from tablename;
```

NOTE: this setting will attempt to do top-level aggregation in the map phase

Sponsored Content



15 Iconic Retail Stores That No Longer Exist
Standard News



Join now to watch Tubelight - a story of one man's unshakeable belief.
Amazon Prime Video



Mansion Global Curates Luxury Real Estate Listings from Around the World
Mansion Global by Dow Jones



18 Feelings Only Extremely Competitive People Will Understand
MensXP



India's Biggest Dance Music Artist, KSHMR Is Touring India Starting Next Month
MensXP.com



Leonardo DiCaprio Buys Malibu Beachfront Home from one of Trump's Cabinet Picks
Mansion Global

Recommended by

NOTE: An aggregation that is NOT top-level is aggregation after GROUP BY

```
hive (mydb)> select count(distinct symbol) from stocks;
237
```

NOTE: The answer will be ZERO if symbol is a partitioned column

```
hive (mydb)> select count(distinct ymd), count(distinct volume) from stocks;
_c0      _c1
```

361	24311
-----	-------

TABLE GENERATING FUNCTIONS(inverse of aggregate functions):

explode, json_tuple, parse_url_tuple, stack

NOTE: aggregate : many rows -> one result

table generating funtions: one column -> many rows/columns

```
hive (mydb)> select explode(subordinates) as sub from employees_no_partition;
sub
Mary Smith
```

OTHER BUILT-IN FUNCTIONS

```
test in(v1, v2, v3, ...) --> returns true if a match is found
length(s)
reverse(s)
concat(s1,s2,s3,...)
concat(separator,s1,s2,s3,...)
substr(s, start, length)
upper(s)/ucase(s),
trim(s), ltrim(s)
regexp_replace(s, regex, replacement) --> regexp_replace(hive, [iv], x) --> hxvx
size(map<k.v>)
siz(array<T>)
cast(<expr> as <type>) --> cast('1' as int)
year, month, day from timestamp
space(n)
repeat(s,n)
split(s, pattern)
instr(str, substr)
in_file(s, filename)
```

LIMIT

```
hive (mydb)> select explode(subordinates) from employees_no_partition limit 1;
```

NESTED SELECT STATEMENT

```
hive> from (select name, salary from employees_no_partition limit 10) e select e.name, e.salary
```

CASE..WHEN..THEN

```
hive> select name, salary,
> case
> when salary < 50000 then 'low'
> when salary >=50000 and salary <70000 then 'middle'
> when salary >=70000 then 'high'
> else 'very high'
> end
> as salary_bracket
> from employees_no_partition;
```

HIVE AVOIDS MAP REDUCE:

```
hive (mydb)> select * from employees_no_partition;
```

NOTE: Hive avoids mapreduce if WHERE clause filters on partition keys

FORCING HIVE TO AVOID MAP REDUCE:

```
hive (mydb)> set hive.exec.mode.local.auto = true;
```

NOTE: No gurantee that Hive won't invoke map reduce

PREDICATE OPERATORS:

```
= (A = B), !=, <, <=, >, >=, IS NULL, IS NOT NULL, LIKE, RLIKE/REGEXP(A RLIKE B --> A= regu
```

NOTE: Use extreme caution when comparing floating point numbers;

Avoid all implicit casts from smaller to wider types;

Avoid floating point numbers for anything involving money

LIKE & RLIKE

```
hive (mydb)> select name from employees_no_partition where address.street like '%Main%';
hive (mydb)> select name from employees_no_partition where address.street rlike '.*(Chicago
```

NOTE: period -> any character

* -> the character appearing immediately left of * can repeat 0 or more times

| -> or

GROUP BY

```
hive (mydb)> select max(price_open) from stocks;
hive (mydb)> select symbol,max(price_open) from stocks;
FAILED: SemanticException [Error 10025]: Line 1:7 Expression not in GROUP BY key 'symbol'
hive (mydb)> select symbol,max(price_open) from stocks GROUP BY symbol limit 5;
hive (mydb)> select symbol,max(price_open) as max from stocks GROUP BY symbol having max >
symbol    max
CACI      49.79
CAF        39.5
CAH        66.9
CAJ        53.22
CAM        42.47
```

JOINS

Hive supports:

- INNER JOIN
- LEFT OUTER JOIN
- RIGHT OUTER JOIN
- FULL OUTER JOIN
- LEFT SEMI JOIN (returns records from the left table if records are found in right table that satisfy ON predicates)
- CARTISIAN PRODUCT JOIN (select * from T1 join T2)
- MAP SIDE JOIN (small table cached in memory) (small table on right side of join)

-> Hive doesn't support OR between predicates in ON clause (AND is supported)

-> You can join more than 2 tables

ORDER BY, SORT BY, DISTRIBUTED BY WITH SORT BY, CLUSTER BY

- ORDER BY -> total ordering
- SORT BY -> orders data within each reducer (local ordering)
- if hive.mapred.mode=strict -> LIMIT needed with ORDER BY
- DISTRIBUTED BY controls how map output is divided among reducers
- CLUSTER BY = DISTRIBUTE BY + SORT BY

-NOTE: Comparison:

ITEM ORDERING # REDUCERS OUTPUT

ORDER BY Global 1 reducer (unacceptable for large data sets) 1 sorted file

SORT BY Local to reducer N reducers (with overlapping data) >=N sorted files

DISTRIBUTE BY No sorting N reducers (non-overlapping data) (not sorted) >=N sorted files

CLUSTER BY [result:global] N reducers (non-overlapping data) (sorted) >=N sorted files

```
hive (mydb)> select s.ymd, s.symbol, s.price_close from stocks s order by s.ymd asc, s.symbol;
hive (mydb)> select s.ymd, s.symbol, s.price_close from stocks s sort by s.ymd asc, s.symbol;
hive (mydb)> select s.ymd, s.symbol, s.price_close from stocks s distribute by s.symbol sort by s.ymd;
hive (mydb)> select s.ymd, s.symbol, s.price_close from stocks s cluster by s.symbol;
```

Example:

```
SORT BY: (each reducer output is ordered; but total order is missing)
r1:      0
0
3
9
r2:      0
0
1
2
ORDER BY: (single output; fully ordered)
r:      0
0
0
0
1
2
3
9
DISTRIBUTE BY: (same keys go into same reducer; no guarantee to be clustered in adjacent positions)
r1:      x1
x2
x1
r2:      x4
x3
CLUSTER BY: (global ordering across multiple reducers)
r1:      x1
```



```
x1
x2
r2:      x3
x4
```

CAST

```
hive (mydb)> select name from employees_no_partition where cast(salary as int) > 100000;
hive (mydb)> select name from employees_no_partition where cast(salary as float) > 100000
hive (mydb)> create external table numbers(number int);
hive (mydb)> load data local inpath 'data/numbers.txt' into numbers;
hive (mydb)> describe numbers;
OK
col_name      data_type      comment
number int
hive (mydb)> select * from numbers;
number
1
2
3
4
5
6
7
8
10
```

TABLESAMPLE-BUCKET

- tablesample (bucket x out of y [on column_name]);
- tablesample (bucket x out of y on rand()) table_alias;
- Rows of the table are bucketed on the column name randomly into y buckets numbered 1 to y
- Rows belonging to bucket x are returned

```
hive (mydb)> select * from numbers TABLESAMPLE(BUCKET 2 OUT OF 5 on rand()) s;
number
7
8
10
```

```
hive (mydb)> select * from numbers TABLESAMPLE(BUCKET 2 OUT OF 5 on rand()) s;
OK
number
hive (mydb)> select * from numbers TABLESAMPLE(BUCKET 2 OUT OF 5 on rand()) s;
OK
number
5
hive> select * from numbers TABLESAMPLE(BUCKET 1 OUT OF 5 on number) s;
OK
1
6
hive> select * from numbers TABLESAMPLE(BUCKET 1 OUT OF 5 on number) s;
OK
1
6
hive> select * from numbers TABLESAMPLE(BUCKET 1 OUT OF 5 on number) s;
OK
1
6
hive> select * from numbers TABLESAMPLE(BUCKET 1 OUT OF 5 on number) s;
OK
5
10
hive> select * from numbers TABLESAMPLE(BUCKET 1 OUT OF 5 on number) s;
OK
5
10
```

BLOCK SAMPLING

→ tablesample(n percent) → n % of the data size (not n% of rows)

```
hive> select * from numbers tablesample(0.1 percent) s;
OK
1
2
3
4
5
6
7
8
10
Time taken: 11.612 seconds
```

–NOTE: Smallest unit of sampling is a single HDFS block; if table

size > block size, all rows are returned;

→ In percentage sampling, by using 'set hive.sample.seednumner=<integer>' you can control the subsets of data sampled

-INPUT PRUNING FOR BUCKET TABLES:

-> Typically, tablesample will scan the entire table; this is not efficient; but if the columns specified in TABLESAMPLE clause match columns in the CLUSTERED BY clause, tablesample queries only scan the required hash partitions of the

```
hive> create table numbers_bucketed(number int) clustered by (number) into 3 buckets;
OK
Time taken: 0.264 seconds

hive> set hive.enforce.bucketing=true;

hive> insert table numbers_bucketed select number from numbers;
FAILED: ParseException line 1:0 cannot recognize input near 'insert' 'table' 'numbers_buck
hive> insert overwrite table numbers_bucketed select number from numbers;
Total MapReduce CPU Time Spent: 20 seconds 600 msec
OK
Time taken: 61.175 seconds

hive> select * from numbers_bucketed;
OK
3
6
1
4
7
10
2
5
8
Time taken: 0.711 second
```

```
[cloudera@localhost ~]$ hadoop fs -cat /user/hive/warehouse/mydb.db/numbers_bucketed/000000
3
6

[cloudera@localhost ~]$ hadoop fs -cat /user/hive/warehouse/mydb.db/numbers_bucketed/000001
1
4
7
10

[cloudera@localhost ~]$ hadoop fs -cat /user/hive/warehouse/mydb.db/numbers_bucketed/000002
2
5
8

hive> select * from numbers_bucketed tablesample(bucket 1 out of 3 on number) s;
OK
3
6

hive> select * from numbers_bucketed tablesample(bucket 2 out of 3 on number) s;
OK
1
4
7
10

hive> select * from numbers_bucketed tablesample(bucket 3 out of 3 on number) s;
OK
2
5
8
```

-UNION ALL (combines 2 or more tables)

-> number and type of columns must match

```
hive> from (from stocks select stocks.symbol, stocks.price_open where stocks.volume < 10000)
> union all
```

```
> from stocks select stocks.symbol, stocks.price_open where stocks.volume <150000)
> unioninput
> insert overwrite directory 'output/union.out' select unioninput.*;

OK

Time taken: 14.794 seconds
```

-NOTE: Above query uses same table for union; same results can be achieved using SELECT & WHERE clause

```
[cloudera@localhost ~]$ hadoop fs -ls output/union.out;
```

Found 1 items

```
-rw-r--r- 3 cloudera supergroup 278413 2015-02-23 14:58 output/union.out/000000_0
```

```
[cloudera@localhost ~]$ hadoop fs -cat output/union.out/000000_0
```

sample_output:

```
CRT30.1
CRT30.7
CRT30.7
CRT31.56
```

VIEWS

-> Materialized views are not supported in Hive

-> Views can be used to reduce query complexity

```
hive> select name, size(subordinates) from employees_no_partition;

OK

John Doe          2
Mary Smith        1
Todd Jones        0
Bill King         0
Boss Man          2
Fred Finance      1
Stacy Accountant  0

Time taken: 20.045 seconds
```

```
hive> create view managers as select * from employees_no_partition where size(subordinates) > 0;

OK
Time taken: 1.185 seconds

hive> show tables;

OK
employees
employees_no_partition
managers
numbers
numbers_bucketed
numbers_clustered
stocks
test

Time taken: 0.484 seconds

hive> select name from managers;

OK
john Doe
Mary Smith
Boss Man
Fred Finance

Time taken: 16.062 seconds
```

```
hive> describe formatted managers;

sample_output:

OK

# col_name      data_type      comment
name            string         None
salary          float          None
```

```
subordinates      array<string>          None
deductions        map<string,float>       None
address           struct<street:string,city:string,
state:string,zip:int>          None
```

Detailed Table Information

```
Database:      mydb
Owner:         cloudera
CreateTime:    Tue Feb 24 00:08:24 CST 2015
Table Type:    VIRTUAL_VIEW
```

View Information

View Original Text: select * from employees_no_partition where size(subordinates) > 0

-NOTE: Table Type: MANAGED_TABLE -> for managed tables

Table Type: EXTERNAL_TABLE -> for external tables

-CREATE TABLES FROM VIEWS:

```
hive> create table tblmanagers like managers;
OK
Time taken: 2.956 seconds
hive> create external table exttblmanagers like managers;
OK
Time taken: 0.757 seconds
```

-ALTER VIEW:

```
hive> alter view managers set tblproperties('created_at'='some_timestamp');
OK
Time taken: 2.435 seconds
```

-DROP VIEW:

```
hive> drop view if exists managers;
```

-INDEX:

```
create table mydb.employees_partitioned(name string,salary float,
subordinates array<string>,deductions map<string, float>,
address struct<street:string, city:string, state:string, zip:int>
partitioned by (country string, state string);
```

[cloudera@localhost ~]\$ cat scripts/create_index.hql;

```
create index employees_partitioned_index on
table mydb.employees_partitioned (country) as 'org.apache.hadoop.hive.ql.index.compact.Com
with deferred rebuild idxproperties
('creator'='me', 'created_at'='some_time')in table mydb.employees_index_tablePARTITIONED BY
(country, name)
comment 'employees_partitioned
by country and name';
```

-???????????????

[cloudera@localhost ~]\$ hive -f 'scripts/create_index.hql';

Logging initialized using configuration in jar:file:/usr/lib/hive/lib/hive-common-0.10.0-cdh4.3.0.jar!/hive-log4j.properties

Hive history file=/tmp/cloudera/hive_job_log_df686526-e57d-49c3-a7e4-7325a6ce2c17_1322001493.txt

FAILED: ParseException line 7:0 missing EOF at 'PARTITIONED' near 'employees_index_table'

-??????????????

-NOTE: Bitmap indexes are commonly used for columns with few distinct values

-> WITH DIFFERED REBUILD, the new index starts empty; at any time, the index can be built the first time or rebuilt using ALTER INDEX statement:

```
alter index employee_index
on table employees
partition(country='US')
rebuild;
```

-> If PARTITION clause is omitted, the index is rebuilt on all partitions

-> show formatted index on employees;

-> drop index if exists employee_index on table employees;

-> AVOID MULTIPLE PASSES OVER THE SAME DATA:

-> **Inefficient:**

```
INSERT OVERWRITE TABLE sales
SELECT * FROM history WHERE action='purchased';

INSERT OVERWRITE TABLE credits
SELECT * FROM history WHERE action='returned';
```

-> **Efficient:**

```
FROM history

INSERT OVERWRITE TABLE sales SELECT * WHERE action='purchased';

INSERT OVERWRITE TABLE credit SELECT * WHERE action='returned';
```

-INDEXES:

-> Hive has limited indexing capabilities

-> There are no keys; but you can build indexes on columns to speed some operations

-> Index data for a table is stored in another table.

```
create index index_name on table table_name (country) as 'org.apache.hadoop.hive.ql.index.c

show index on table_name;

drop index index_name on table_name;
```

-EXPLAIN:

```
hive> explain select symbol, count(symbol) from stocks group by symbol;
OK
```

ABSTRACT SYNTAX TREE:

```
(TOK_QUERY (TOK_FROM (TOK_TABREF (TOK_TABNAME stocks))) (TOK_INSERT (TOK_DESTINATION (TOK_C
```

STAGE DEPENDENCIES:

Stage-1 is a root stage

Stage-0 is a root stage

STAGE PLANS:

```

Stage: Stage-1

Map Reduce

Alias -> Map Operator Tree: stocks

TableScan

alias: stocks

Select Operator

expressions:

expr: symbol

type: string

outputColumnNames: symbol

Group By Operator

aggregations:

expr: count(symbol)

bucketGroup: false

keys:

expr: symbol

type: string

mode: hash

outputColumnNames: _col0, _col1

Reduce Output Operator

```

–NOTE: You can also use:

explain EXTENDED select symbol, count(symbol) from stocks group by symbol;

—>TUNING

- Use [EXTENDED] EXPLAIN
- Tune LIMIT clause (use set)
- Tune JOIN clause (use map-side join on small datasets)
- use Local Mode on small datasets (use set)
- Parallel Execution (execute jobs in parallel) (use set)
- Strict Mode (use set hive.mapred.mode=strict). Enforces the following:
 - enforces WHERE clause on partitioned tables
 - enforces LIMIT clause with ORDER BY
 - enforces ON clause with cartesian product (JOIN)
- Tuning number of reducers:
 - set hive.exec.reducers.bytes.per.reducer=750,000,000
 - set hive.exec.reducers.max (max reducer count)
 - reduce.tasks (reducer count)
- JVM Reuse (use set mapred.job.reuse.jvm.num.tasks=10)
- Use Indexes
- Tune dynamic partition:
 - exec.dynamic.partition.mode=strict (at least one static partition is needed)
 - exec.max.dynamic.partitions
 - exec.max.dynamic.partitions.pernode
- Speculative Execution (Hadoop launches multiple instances of map/reduce jobs):
 - map.tasks.speculative.execution=true;
 - reduce.tasks.speculative.execution=true;
 - mapred.reduce.tasks.speculative.execution=true;
- Use single mapreduce job to combine multiple GROUP BY in a query (common group by key is rdd) (use set)
- Use Virtual columns (see example below)

```

hive> select INPUT__FILE__NAME, BLOCK__OFFSET__INSIDE__FILE, symbol
from stocks where symbol like '%CLIX%' limit 5;

hdfs://localhost.localdomain:8020/user/cloudera/myhive/NYSE_daily      0
hdfs://localhost.localdomain:8020/user/cloudera/myhive/NYSE_daily      57

```

```

hdfs://localhost.localdomain:8020/user/cloudera/myhive/NYSE_daily      114
hdfs://localhost.localdomain:8020/user/cloudera/myhive/NYSE_daily      171
hdfs://localhost.localdomain:8020/user/cloudera/myhive/NYSE_daily      228
Time taken: 17.629 seconds

```

NOTE: INPUT__FILE__NAME & BLOCK__OFFSET__INSIDE__FILE needs to be in upper case; seperator is 2 dashes(__)

SCRIPTS:

```
[cloudera@localhost ~]$ cat scripts/create_table.hql
```

```

CREATE TABLE IF NOT EXISTS mydb.employees (
  name STRING COMMENT 'Employee name',
  salary FLOAT COMMENT 'Employee salary',
  subordinates ARRAY<STRING> COMMENT 'Names of subordinates',
  deductions MAP<STRING, FLOAT> COMMENT 'Keys are deductions names, values are percentages',
  address STRUCT<street:STRING, city:STRING, state:STRING, zip:INT> COMMENT 'Home address')
COMMENT 'Description of the table'
LOCATION '/user/hive/warehouse/'
TBLPROPERTIES ('creator'='me', 'created_at'='2015-02-19 9:55:00');

```

```
[cloudera@localhost ~]$ hive -f 'scripts/create_table.hql'
```

OTHERS:

-CHANGING HIVE PROMPT:

```

hive> set hive.cli.print.current.db=true;
hive (default)>

```

-PRINT column headers;

```
hive> set hive.cli.print.header=true;
```

-HIVE ONESHOT COMMANDS:

```
[cloudera@localhost ~]$ hive -e 'select * from records LIMIT 5';
```

```
[cloudera@localhost ~]$ hive -e 'select * from mydb.stocks limit 5;'
```

-EXECUTE UNIX COMMANDS:

```

hive (test)> !pwd;
hive (test)> /home/cloudera

```

-INSTALLATION LOCATION:

```
[cloudera@localhost ~]$ ls /usr/lib
```

-HIVE INSTALLATION LOCATION

```
[cloudera@localhost ~]$ ls /usr/lib/hive/lib
```

```
cloudera@localhost ~]$ ls -l /usr/lib/hive/bin
```

total 28

```
-rwxr-xr-x. 1 root root 881 May 27 2013 beeline
```

```
drwxr-xr-x. 3 root root 4096 Jun 18 2013 ext
```

```
-rwxr-xr-x. 1 root root 5781 May 27 2013 hive
```

```
-rwxr-xr-x. 1 root root 1900 May 27 2013 hive-config.sh
```

```
-rwxr-xr-x. 1 root root 885 May 27 2013 hiveserver2
-rwxr-xr-x. 1 root root 832 May 27 2013 metatool
[cloudera@localhost ~]$
```

–HIVE CONFIGURATION:

```
[cloudera@localhost ~]$ ls /usr/lib/hive/conf
```

–HIVE LOG:

[NOTE: cat /tmp/cloudera/hive.log displays the log contents]

–HIVE LOGGING:

Open following file in text editor and you can change logging level:

```
/usr/lib/hive/lib/hive-common-0.10.0-cdh4.3.0.jar/hive-log4j-properties
```

```
[cloudera@localhost ~]$ hive -help
```

usage: hive

-d, –define <key=value> Variable substitution to apply to hive

commands. e.g. -d A=B or –define A=B

–database <databasename> Specify the database to use

-e <quoted-query-string> SQL from command line

-f <filename> SQL from files

-H, –help Print help information

-h <hostname> connecting to Hive Server on remote host

–hiveconf <property=value> Use value for given property

–hivevar <key=value> Variable substitution to apply to hive

commands. e.g. –hivevar A=B

-i <filename> Initialization SQL file

-p <port> connecting to Hive Server on port number

-S, –silent Silent mode in interactive shell

-v, –verbose Verbose mode (echo executed SQL to the console)

–HIVE SERVICES:

```
[cloudera@localhost ~]$ hive –service help
```

[NOTE: cli : command line interface (default)]

hiveserver: runs hive as a server (exposes a thrift service)

hwi: hive web interface

jar: similar to hadoop jar

metastore: runs metastore as a standalone process]

[NOTE: Thrift client: lets you run hive commands from different programming languages]

```
[cloudera@localhost ~]$ hive –help
```

Usage: ./hive <parameters> –service serviceName <service parameters>

Service List: beeline cli help hiveserver2 hiveserver hwi jar lineage metastore metatool rcfilecat

Parameters parsed:

–auxpath : Auxillary jars

–config : Hive configuration directory

–service : Starts specific service/component. cli is default

Parameters used:

HADOOP_HOME or HADOOP_PREFIX : Hadoop install directory

HIVE_OPT : Hive options

For help on a particular service:

./hive –service serviceName –help

Debug help: ./hive –debug –help

–CLI SERVICE

```
[cloudera@localhost ~]$ hive –help –service cli
```

usage: hive

–d, –define <key=value> Variable substitution to apply to hive commands. e.g. –d A=B or –define A=B

–database <databasename> Specify the database to use

–e <quoted-query-string> SQL from command line

–f <filename> SQL from files

–H, –help Print help information

–h <hostname> connecting to Hive Server on remote host

–hiveconf <property=value> Use value for given property

–hivevar <key=value> Variable substitution to apply to hive commands. e.g. –hivevar A=B

–i <filename> Initialization SQL file

–p <port> connecting to Hive Server on port number

–S, –silent Silent mode in interactive shell

–v, –verbose Verbose mode (echo executed SQL to the console)

– HIVE NAMESPACES:

hivevar

hiveconf

system

env

–SELECT in NON-INTERACTIVE MODE:

```
[cloudera@localhost ~]$ hive –e ‘select * from records’;
```

Logging initialized using configuration in

```
[cloudera@localhost ~]$ hive –S –e “set” | grep warehouse
```

hive.metastore.warehouse.dir=/user/hive/warehouse

hive.warehouse.subdir.inherit.perms=true

–HIVE FILES:

```
[cloudera@localhost ~]$ cat scripts/select.q
```

select * from records;

–CHANGE TO ROOT USER

```
[cloudera@localhost ~]$ su root
```

Password: →cloudera

```
[root@localhost cloudera]#
```

–EXECUTE HADOOP COMMANDS:

```
hive (test)> dfs -ls output;
```

–LIST ALL HADOOP COMMANDS:

```
hive> dfs -help;
```

–COMMENTS IN HIVE:

–this is a comment

WORK:

```
hadoop1@ubuntu:~/scripts$ cat nyse_daily.hql
```

```
create external table if not exists
mydb.nyse_daily (exchange string,symbol string,
current_date date,open float,
high float,low float,
close float,volume int,
adj_close float)row format delimited
fields terminated by '\t';
```

```
hadoop1@ubuntu:~/scripts$ hive -f 'nyse_daily.hql'
```

```
hive (mydb)> load data inpath '/home/hadoop1/data/NYSE_daily' overwrite into table nyse_dai
```

FAILED: SemanticException Line 1:17 Invalid path "/home/hadoop1/data/NYSE_daily": No files matching path
hdfs://localhost:9000/home/hadoop1/data/NYSE_daily

```
hadoop1@ubuntu:~/data$ hadoop fs -put NYSE* data
```

```
hive (mydb)> load data inpath 'data/NYSE_daily' overwrite into table nyse_daily;
hive (mydb)> select * from nyse_daily limit 10;
```

Share this:

 Tweet



About Siva

Senior Hadoop developer with 4 years of experience in designing and architecture solutions for the Big Data domain and has been involved with several complex engagements. Technical strengths include Hadoop, YARN, Mapreduce, Hive, Sqoop, Flume, Pig, HBase, Phoenix, Oozie, Falcon, Kafka, Storm, Spark, MySQL and Java.

[View all posts by Siva →](#)



Leave a comment

Your email address will not be published. Required fields are marked *

b *i* link **b-quote** ~~del~~ **ins** **img** **ul** **ol** **li** **code** **more** **close tags** **crayon** **¶**

Name *

Email *

Website

2 thoughts on “Hive Functions Examples”

**Suresh Perumal***June 6, 2015 at 5:31 pm*

Dear Siva

Surely this posting will be useful for [hadoop](http://itcoordinatestraining.com) developer . Clearly explained with examples. thank u.

with Regards

Suresh Perumal

ITCoordinates Mylapore

**Ms Revathi***July 1, 2016 at 8:28 pm*

such an awesome site this is...Loved it..very useful and helpful examples.

Post navigation

[← Pig Functions Examples](#)[100 Hadoop Certification Dump Questions →](#)[Review Comments](#)



Course content is well structured. I like Siva's explanation of topics using slide decks & virtual machine (CDH cluster) at the same time, this will help audience to learn not only theory behind a topic but also practical aspect of it. Overall, I would recommend this course.

Kumar
Big Data Developer
Hadoop&Aug/2016
September 26, 2016



Sponsored Content




5 Multi-bagger Stocks for the next 5 years
www.5paisa.com



The BMW X1. Celebrate as you beat the price hike. Know more.
BMW India



Eating is the new way to stay fit!
Saffola Fit Foodie



12 Women Who Broke The Internet With Their Hotness
MensXP-Special Features

Recommended by



Recent Comments

›WENJIE LI on Creating Custom Hadoop Writable Data Type
›Brian on Avro Serializing and Deserializing Example – Java API
›chandan on Hive Authorization Models and Hive Security
›venkat on Hive Database Commands
›Siva on Avro Serializing and Deserializing Example – Java API

Hadooptutorial.info

Hadooptutorial.info

Let's get Social :



© 2017 Hadoop Online Tutorials · Designed by
Press Customizr ·

[Back to top](#)

[Big Data](#)

[Hard Drive](#)

[Online Data](#)




[ECO Data](#)

[Hitachi-LG Data](#)

ads by Yahoo!

Custom Search

- [Home](#)
- [Big Data](#)
- [Hadoop](#)
- [Map Reduce](#)
- [Hive](#)
- [Pig](#)
- [HBase](#)
- [Flume](#)
- [Interview Questions](#)
- [Miscellaneous](#)

-  [Twitter](#)
-  [Facebook](#)
-  [Google+](#)

»